



# Support Vector Machines


Danushka Bollegala  
Lecture 12

# Linear Separability

- Consider binary classification of two dimensional feature vectors
  - e.g. features = {good, bad}
  - classes = {positiveSentiment, negativeSentiment}
- If we can find a straight line that can separate all positive instances (reviews) from all negative instances (reviews) then we call such a dataset to be *linearly separable*







# Higher Dimensions

- Reviews contain more than two features (words)
- In  $N$ -dimensional space, we must find  $(n-1)$  dimensional hyperplane that separates the two classes (if they are linearly separable)
- $n=2$  (two dimensional feature space), we had straight lines ( $n=1$  dimensional hyperplanes)
- Hyperplane that separates the two classes might not be unique (as we saw in our previous example)

# Large Margin Classifiers


- Find **two** hyperplanes that separates the positive class and the negative class
- Try to maximise the minimum separation (distance) between the two hyperplanes
- The distance between the hyperplanes is called the **margin**
- Maximising the margin minimises the **risk** of misclassifying an instance at test time
- reduces overfitting

# Support Vector Machines

- Support Vector Machines (SVMs) are one of the many large margin classification methods
- Uses a constrained convex optimisation method
- Can handle non-linear separable datasets using
  - slack variables
  - kernel functions



# SVM Idea




# Distance to a straight line

- Given a straight line  $l: ax+by+c=0$  show the perpendicular distance  $d$  to  $l$  from a point  $(\alpha, \beta)$  is

## Home Work 1

$$d = \frac{a\alpha + b\beta + c}{\sqrt{a^2 + b^2}}$$



# Distance to a hyperplane

- A hyperplane can be expressed as the inner-product between a weight vector (coefficients) and a feature vector (variables corresponding to the dimensions)
- $\mathbf{w}^\top \mathbf{x} = 0$
- Then the distance to this hyperplane from a point  $p$ , given by the vector  $\mathbf{p}$  can be computed as

$$\frac{\mathbf{w}^\top \mathbf{p}}{\|\mathbf{w}\|}$$

- where  $\|\mathbf{w}\|$  is the norm (L2 length) of the vector  $\mathbf{w}$
- Observe that this formula reduces to the one we derived in the two-dimensional case in the previous slide

# SVM background

- Let us assume we are given a training dataset  $(t_n, x_n)$  of  $n=1, \dots, N$  instances
  - target labels  $t_n = \{-1, +1\}$  for binary classification
- The feature vector for the instance  $x$  is represented by  $\phi(x)$
- Our classification decision of  $x$  is made according to the score  $y(x)$  given by
  - $y(x) = \mathbf{w}^T \phi(x) + b$
- Here,  $\mathbf{w}$  is the weight vector and  $b$  is the bias (scalar) term that adjust any fixed bias from the 0 threshold
  - If  $y(x) > 0$  then we classify  $x$  to be positive and
  - otherwise negative

# SVM Derivation

- If a point (instance) is correctly classified by the hyperplane then
  - $t_n y(\mathbf{x}_n) > 0$
- The distance from a correctly classified point to the hyperplane is given by

$$\frac{t_n y(\mathbf{x}_n)}{\|\mathbf{w}\|} = \frac{t_n (\mathbf{w}^T \phi(\mathbf{x}_n) + b)}{\|\mathbf{w}\|}.$$

# SVM Derivation

- We need to find the weight vector  $\mathbf{w}$  and bias term  $b$  such that this margin is maximised for all the training instances in our train dataset

$$\arg \max_{\mathbf{w}, b} \left\{ \frac{1}{\|\mathbf{w}\|} \min_n \left[ t_n \left( \mathbf{w}^T \phi(\mathbf{x}_n) + b \right) \right] \right\}$$

This is a difficult optimisation problem involving min-max. Moreover, it is scale-invariant meaning that by setting  $\mathbf{w} \rightarrow k\mathbf{w}$  and  $b \rightarrow kb$  the term inside min does not change!

# Simplification!

- Scale the parameters such that a point on the decision hyperplane satisfies

$$t_n (\mathbf{w}^T \phi(\mathbf{x}_n) + b) = 1$$

- All correctly classified data points will then satisfy

$$t_n (\mathbf{w}^T \phi(\mathbf{x}_n) + b) \geq 1, \quad n = 1, \dots, N.$$

- This is called the *canonical representation* of the decision hyperplane

# SVM Derivation 3

- Now the margin becomes

$$\frac{t_n y(x_n)}{\|\mathbf{w}\|} = \frac{t_n (\mathbf{w}^\top \phi(x_n) + b)}{\|\mathbf{w}\|} = \frac{1}{\|\mathbf{w}\|}$$

- Great!
- Now our final objective becomes to find  $\mathbf{w}$  and  $b$  such that we **maximise the margin** subjected to the set of **constraints that ensures our train data instances are correctly classified**
- Maximising margin = minimising the norm  $\|\mathbf{w}\|$



# SVM Optimisation Problem

- Find  $\mathbf{w}$  and  $b$  such that

- minimise

$$\min \frac{1}{2} ||w||^2$$

- subjected to

$$t_n \left( \mathbf{w}^T \phi(\mathbf{x}_n) + b \right) \geq 1, \quad n = 1, \dots, N.$$

# Constrained Optimisation

- Find  $x$  that minimises  $f(x)$ 
  - unconstrained optimisation
- Find  $x$  that minimises  $f(x)$  subjected to  $g(x) = 0$ 
  - constrained optimisation

# unconstrained vs. constrained

- minimise  $f(x,y) = x^2 + y^2$
- such that  $g(x,y) = y - x - 1 = 0$


$$f(x,y) = x^2 + (x+1)^2 = 2x^2 + 2x + 1$$

$$\frac{\partial f}{\partial x} = 4x + 2 = 0$$

$$x = -1/2$$

$$y = x + 1 = 1/2$$

$$\min f(x,y) = (-0.5)^2 + (0.5)^2 = 0.5$$



# Lagrange Multipliers

- Problem:
  - Minimise  $f(x)$  subjected to  $g(x) \geq 0$
- Lagrangian function for the problem becomes
  - $L(x, \lambda) = f(x) - \lambda g(x)$
  - $\lambda \geq 0$  is called the Lagrange variable
- Procedure
  - Compute  $x$  and  $\lambda$  by solving

$$\frac{\partial L(x, \lambda)}{\partial \lambda} = 0$$
$$\frac{\partial L(x, \lambda)}{\partial w} = 0$$

# Idea

Minimising a two variable function  $f(x,y)$  w.r.t.  $x$  and  $y$  means that we are drawing the contours for  $f(x,y)$ .




figure from Wikipedia

Minimising while satisfying  $g(x,y) = c$  happens when the two curves touch each other.

At this point the two gradients must be parallel and in opposite directions

# Home Work

- Use Lagrangian multiplier method to solve the optimisation problem in slide 19

# Back to SVM Derivation

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N a_n \{t_n(\mathbf{w}^T \phi(\mathbf{x}_n) + b) - 1\}$$

$$\frac{\partial L(\mathbf{w}, b, \mathbf{a})}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)$$

$$\frac{\partial L(\mathbf{w}, b, \mathbf{a})}{\partial b} = 0 \Rightarrow 0 = \sum_{n=1}^N a_n t_n.$$

# SVM

- Plugging these back to the Lagrangian function we get

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\mathbf{x}_n, \mathbf{x}_m)$$

- Which must be solved subjected to constraints

~

$$a_n \geq 0, \quad n = 1, \dots, N,$$

$$\sum_{n=1}^N a_n t_n = 0.$$



# Observations

- We must find Lagrange multipliers  $a_n$  (collectively denoted by the vector  $\mathbf{a}$ ) such that  $L(\mathbf{a})$  is minimised.
- We have the inner-product between two instances  $x_n$  and  $x_m$  appearing in the objective function
  - $k(x_n, x_m) = \phi(x_n)^T \phi(x_m)$
  - Only the inner products matter. We do not need the explicit form of feature vectors  $\phi(x)$
  - Can be *kernalised* using numerous kernel functions to overcome the non-linear separability issue.

# Observations

- Note that if the Lagrange multiplier  $a_n = 0$ , then the  $n$ -th instance has no effect on the objective function  $L$
- The instances that correspond to non-zero Lagrange multipliers are the ones that we need to store in our final model
- Support Vectors
  - The instances that appear on top of the decision hyperplanes and determine its shape

# Classification with SVMs

- During test time, to classify a test instance  $\mathbf{x}$ , we simply compute the inner-product between  $\mathbf{x}$  and each of the support vectors  $\mathbf{x}_n$

$$y(\mathbf{x}) = \sum_{n=1}^N a_n t_n k(\mathbf{x}, \mathbf{x}_n) + b.$$

We still do not need the explicit representation of  $\mathbf{x}$  or  $\mathbf{x}_n$  and can work with the values returned by the kernel function

# Home Work

- Using the decision function  $y(x) = \mathbf{w}^T \phi(x) + b$  and the result we obtained for  $\mathbf{w}$  in slide 23, derive the classification function for SVMs in the kernel form as shown in slide 27.

# Kernel Functions

- Linear Kernel

$$k(\mathbf{x}_n, \mathbf{x}_m) = \mathbf{x}_n^\top \mathbf{x}_m$$

- Does not use any transformations

- Polynomial Kernel

$$k(\mathbf{x}_n, \mathbf{x}_m) = (\mathbf{x}_n^\top \mathbf{x}_m + c)^d$$

- Quadratic (d=2), and Cubic (q=3) are widely used.
- Can account for the combinations of features such as bigrams in text mining tasks

- Sigmoid Kernel

$$k(\mathbf{x}_n, \mathbf{x}_m) = \tanh(\mathbf{x}_n^\top \mathbf{x}_m + c)$$

- Exponential Radial Basis Function (RBF) Kernel


$$k(\mathbf{x}_n, \mathbf{x}_m) = \exp\left(-\frac{\|\mathbf{x}_n - \mathbf{x}_m\|^2}{2\sigma^2}\right)$$

- Subsumes all possible kernel functions

# Slack variables

- Sometimes it is easy to *shift* some of the training instances (especially around the decision hyperplane) so that the dataset becomes linearly separable
- Doing this too much will change the train data significantly and we will not learn the concept expressed by our train data
- Try to minimise the amount of shifting we do for train instances to make the problem linearly separable
- Each train instance is associated with a slack variable that is set to a non-zero value such that the corresponding training instance is moved sufficiently to the correct side of the decision hyperplane

# SVMs and slack variables



$\xi_n > 1$  misclassification

$$t_n y(\mathbf{x}_n) \geq 1 - \xi_n, \quad n = 1, \dots, N$$

## slacked version of the constraint

$$C \sum_{n=1}^N \xi_n + \frac{1}{2} \|\mathbf{w}\|^2$$

# objective function

C: cost-parameter

Higher values of  $C$  impose heavier penalties of slacking, whereas smaller  $C$  values will change the train data significantly. In practice use cross-validation to set  $C$ .

# SVM slack version

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N a_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} - \sum_{n=1}^N \mu_n \xi_n$$

$$\frac{\partial L}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)$$

$$\frac{\partial L}{\partial b} = 0 \Rightarrow \sum_{n=1}^N a_n t_n = 0$$

$$\frac{\partial L}{\partial \xi_n} = 0 \Rightarrow a_n = C - \mu_n.$$

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\mathbf{x}_n, \mathbf{x}_m)$$

Same Lagrangian as before!



# SVM Implementations

- LIBSVM
  - <http://www.csie.ntu.edu.tw/~cjlin/libsvm/>
  - Available in a large number of programming languages
- SVM Light
  - <http://svmlight.joachims.org/>
  - can do ranking SVMs

# Home Work

- Use LIBSVM to train a binary sentiment classifier using the train data provided in the Assignment 1
- Compare the performance with the perceptron classifier that you implemented