

Visualization Theory

Multidimensional Scaling

Danushka Bollegala



Objective

- Assume that we have a dataset that we would like to visualize
- Problem settings
 - Our data are **multi-dimensional**
 - Apply a dimensionality reduction method such as PCA.
 - Our data points are **scalar** values (e.g. distances/similarities between some points)
 - Apply **Multi-dimensional scaling (MDS)**

Multidimensional Scaling (MDS)

- Assume that we have distance/similarity matrix ($n \times n$) for n data points.
- Examples
 - We asked human subjects to express their preferences
 - We extracted clickthrough data from a search engine
- How can we construct a set of n feature vectors for the data points, given their similarity matrix?
- Problem tackled by MDS

Definitions

- Assume that we are given a distance matrix $\mathbf{D} \in \mathbb{R}^{n \times n}$.
- We are interested in finding a set of p dimensional vectors arranged in rows of a matrix $\mathbf{X} \in \mathbb{R}^{n \times p}$.
- We first consider computing a matrix $\mathbf{B} = \mathbf{X}\mathbf{X}^T$ using \mathbf{D} .
- If we can find \mathbf{B} then we can easily compute \mathbf{X} as we shall see later.

Step 1: Express D in terms of B

- Let us denote the (r,s) element of **B** by b_{rs}
- The squared Euclidean distance, d_{rs} , between the data point r and the data point s is given by

$$d_{rs}^2 = ||\mathbf{x}_r - \mathbf{x}_s||_2^2$$

$$d_{rs}^2 = \sum_{i=1}^p (x_{ri} - x_{si})^2$$

$$d_{rs}^2 = \sum_{i=1}^p x_{ri}^2 - 2 \sum_{i=1}^p x_{ri} x_{si} + \sum_{i=1}^p x_{si}^2$$

$$d_{rs}^2 = b_{rr} - 2b_{rs} + b_{ss} \quad [1]$$

Standardization of \mathbf{X}

- It does not matter how we select the origin of the co-ordinate space.
- Let us assume that all features in \mathbf{X} are standardized such that their mean value is zero.
- This means the sum of each column of \mathbf{X} is zero.

$$\sum_{r=1}^n x_{ri} = 0 \quad \forall i = 1, \dots, p \quad [2]$$

- By standardization we reduce the freedom of the matrix \mathbf{X} , thereby constraining it.

Proof: Sum of a row r in \mathbf{B} is zero

- If \mathbf{X} is standardized, then the sum of any row r in \mathbf{B} becomes zero.

Proof:

$$\begin{aligned}\sum_{s=1}^n b_{rs} &= \sum_{s=1}^n \sum_{i=1}^p x_{ri} x_{si} \\ &= \sum_{i=1}^p \sum_{s=1}^n x_{ri} x_{si} \\ &= \sum_{i=1}^p x_{ri} \sum_{s=1}^n x_{si} \\ &= 0 \quad \blacksquare \quad [3]\end{aligned}$$

Proof: Sum of a column s in \mathbf{B} is zero

- If \mathbf{X} is standardized, then the sum of any column s in \mathbf{B} becomes zero.

Proof:

$$\begin{aligned}\sum_{r=1}^n b_{rs} &= \sum_{r=1}^n \sum_{i=1}^p x_{ri} x_{si} \\ &= \sum_{i=1}^p \sum_{r=1}^n x_{ri} x_{si} \\ &= \sum_{i=1}^p x_{si} \sum_{r=1}^n x_{ri} \\ &= 0 \quad \blacksquare \quad [4]\end{aligned}$$

Lets sum-up the squared distances

$$d_{rs}^2 = b_{rr} - 2b_{rs} + b_{ss}$$

$$\sum_{r=1}^n d_{rs}^2 = \sum_{r=1}^n b_{rr} - 2 \sum_{r=1}^n b_{rs} + \sum_{r=1}^n b_{ss} = \text{tr}(\mathbf{B}) + nb_{ss} \quad [5]$$

sum of diagonal elements of B. This is called the trace of a matrix and is denoted by tr.

This term is zero from [4].

We are adding n times b_{ss} , which has nothing to do with r.

Lets sum-up the squared distances

$$d_{rs}^2 = b_{rr} - 2b_{rs} + b_{ss}$$

$$\sum_{s=1}^n d_{rs}^2 = \sum_{s=1}^n b_{rr} - 2 \sum_{s=1}^n b_{rs} + \sum_{s=1}^n b_{ss} = \text{tr}(\mathbf{B}) + nb_{rr} \quad [6]$$

We are adding
n times b_{rr} , which has
nothing to do with r.

This term is zero
from [3].

sum of diagonal
elements of B.
This is called the
trace of a matrix
and is denoted by
tr.

Lets sum-up the squared distances

$$\begin{aligned}\sum_{r=1} \sum_{s=1}^n d_{rs}^2 &= \sum_{r=1} \text{tr}(\mathbf{B}) + n \sum_{r=1} b_{rr} \\ &= n \text{tr}(\mathbf{B}) + n \text{tr}(\mathbf{B}) \\ &= 2n \text{tr}(\mathbf{B})\end{aligned}$$

Taking the sum over r in [6].

[7]

Quiz: Derive [7] by taking the sum over s in [5].

Express d_{rs} in terms of b_{rs}

- Substitute for b_{rr} and b_{ss} in [1] using [5] and [6]

$$d_{rs}^2 = b_{rr} - 2b_{rs} + b_{ss}$$
$$= \frac{1}{n} \left(\sum_{r=1}^n d_{rs}^2 - \text{tr}(\mathbf{B}) \right) - 2b_{rs} + \frac{1}{n} \left(\sum_{s=1}^n d_{rs}^2 - \text{tr}(\mathbf{B}) \right)$$

- Further substitute for $\text{tr}(\mathbf{B})$ from [7]

$$d_{rs}^2 = b_{rr} - 2b_{rs} + b_{ss}$$
$$= \frac{1}{n} \left(\sum_{r=1}^n d_{rs}^2 - \frac{1}{2n} \sum_{r=1}^n \sum_{s=1}^n d_{rs}^2 \right) - 2b_{rs} + \frac{1}{n} \left(\sum_{s=1}^n d_{rs}^2 - \frac{1}{2n} \sum_{r=1}^n \sum_{s=1}^n d_{rs}^2 \right)$$

[8]

Computing **B**

- From [8] we have,

$$b_{rs} = -\frac{1}{2} \left[d_{rs}^2 - \frac{1}{n} \sum_{r=1}^n d_{rs}^2 - \frac{1}{n} \sum_{s=1}^n d_{rs}^2 + \frac{1}{n^2} \sum_{r=1}^n \sum_{s=1}^n d_{rs}^2 \right] \quad [9]$$

Using [9] we can compute all elements in **B**. The right hand side of [9] contains only the distances d_{rs} , which are given to us.

So at last we have the matrix **B** that we needed.

Also note that **B** is a symmetric matrix because the distance matrix **D** is a symmetric matrix.

Take a deep breath!

Are we done?

- Not quite! We must compute \mathbf{X} from $\mathbf{B} = \mathbf{X}\mathbf{X}^\top$
- Easy... Perform eigenvalue decomposition on \mathbf{B}
 - $\mathbf{B} = \mathbf{U}\mathbf{G}\mathbf{U}^\top$
 - where, \mathbf{U} is an orthogonal matrix $\mathbf{U}\mathbf{U}^\top = \mathbf{U}^\top\mathbf{U} = \mathbf{I}$
 - \mathbf{G} is a diagonal matrix containing the eigenvalues of \mathbf{B}
 - Let $\mathbf{X} = \mathbf{U}\mathbf{G}^{1/2}$, then we have

$$(\mathbf{U}\mathbf{G}^{1/2})(\mathbf{U}\mathbf{G}^{1/2})^\top = \mathbf{U}\mathbf{G}^{1/2}\mathbf{G}^{1/2}\mathbf{U}^\top = \mathbf{U}\mathbf{G}\mathbf{U}^\top = \mathbf{B}$$

MDS Algorithm

- INPUT
 - Distance matrix **D**, dimensionality p .
- Output
 - Matrix **X** where each row represents a feature vector for a data point
- Compute **B** using
$$b_{rs} = -\frac{1}{2} \left[d_{rs}^2 - \frac{1}{n} \sum_{r=1}^n d_{rs}^2 - \frac{1}{n} \sum_{s=1}^n d_{rs}^2 + \frac{1}{n^2} \sum_{r=1}^n \sum_{s=1}^n d_{rs}^2 \right]$$
- Perform eigenvalue decomposition **B** = **UGU**^T on **B** to obtain **U** and **G**. Select the top p eigenvalues/vectors.
- Return **X** = **UG**^{1/2}

Things to note...

- Your distance matrix **D** might have negative eigenvalues
 - which means it is not Euclidean. You cannot exactly reproduce it using **X**. But this is fine as long as you have some large positive eigenvalues
- Matrix **B** might have negative eigenvalues
 - You get complex values when you compute their square root
 - If you have some high positive values (at least 2), then you can still plot in the 2D space
- The point is to construct vectors **X** that produce the distances in **D** as close as possible
 - We can measure the reproduction error using root mean square error (RMSE) [see classifier evaluations lecture]


```

import numpy
import matplotlib.pyplot as plt

def bval(D, r, s):
    n = D.shape[0]
    total_r = numpy.sum(D[:,s] ** 2)
    total_s = numpy.sum(D[r,:] ** 2)
    total = numpy.sum(D ** 2)
    val = (D[r,s] ** 2) - (float(total_r) / float(n)) - (float(total_s) / float(n)) + (float(total) / float(n * n))
    return -0.5 * val


def main():
    n = 50; m = 2
    Y = numpy.random.rand(n, m)
    D = numpy.zeros((n, n), dtype=complex)
    for i in range(0, n):
        for j in range(0, n):
            D[i, j] = numpy.linalg.norm(Y[i] - Y[j])
    B = numpy.zeros((n, n), dtype=complex)
    for i in range(0, n):
        for j in range(0, n):
            B[i,j] = bval(D, i, j)
    print "\nB matix"
    print B
    g, U = numpy.linalg.eig(B)
    idx = g.argsort()[::-1]
    g = g[idx]
    U = U[:,idx]
    print "Eigenvalues =", g
    G = numpy.diag(numpy.sqrt(g))
    X = numpy.dot(U.T, G)
    print "\nMatrix X"
    print X
    error = 0.0
    total = 0
    for i in range(0, n):
        for j in range(i+1, n):
            error += (numpy.linalg.norm(X[i] - X[j]) - D[i, j]) ** 2
            total += 1
    print "RMSE =", numpy.sqrt(error / float(total))
    plt.plot(X[:,0], X[:,1], 'bo', markersize=14)
    plt.plot(Y[:,0], Y[:,1], 'ro', markersize=14)
    plt.show()
    pass

if __name__ == '__main__':
    main()


```

Example 1: 2D data

- Let us consider a set 50 of 2D random feature vectors.
- We compute a distance matrix **D** from these vectors and perform MDS (p=2) on **D**.
- We plot the original vectors (red) and the vectors produced by MDS (blue) in 2D.
- Observe the correlation between RMSE and the MDS results.




RMSE = 0.299, m = 2



RMSE = 0.472, m = 2

Example 2: High-dimensional Data

- Let us consider a set of 50 random vectors in 1000 dimensional space.
- We compute a distance matrix \mathbf{D} from these vectors and perform MDS ($p=2$) on \mathbf{D} .
- We plot the original vectors and the vectors produced by MDS in 2D space
- It is not possible to see any difference among original data points (shown in red) in the 2D space, if we only consider their first two dimensions.
- However, MDS provides a much better separation among the data points (shown in blue)



m = dimensionality of
the actual space
 $m = 1000$
RMSE = 0.3965